

## Neural Code Generation Models with Programming Language Knowledge

Yingfei Xiong Peking University

## Training Code LLM



More parameters, more effective, more expensive

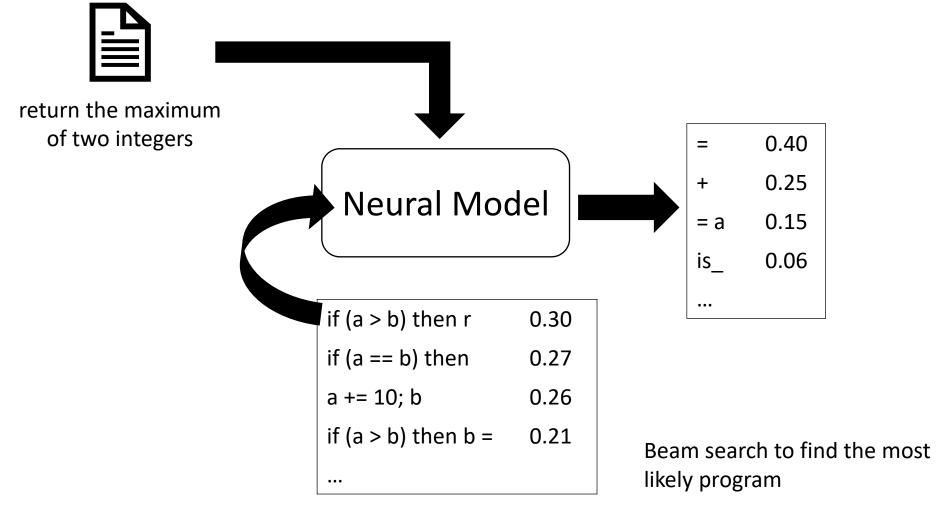


Large model costs millions and even billions to train

• Can we reduce the needed parameters, or improve performance without using more parameters?

## Neural Code Generation: Predicting BPE tokens one by one





## Problem: Ignoring PL knowledges



 Program languages have knowledge such as grammar, type system, and semantics, ignoring them making learning more difficult.

Syntactic constraint: ()+5( illegal

Type constraint: 1+true illegal

Semantic constraint:

use without initialization illegal

Functions between token sequences

Functions between compilable programs

Hypothesis space (PAC learning)

## Problem: Ignoring PL knowledges



Correct code generation requires PL knowledges

assignment is common and should be used

NN that does not know types

all parameters are Boolean so 'if' is more likely

NN that knows types

## Problem: Ignoring PL knowledges



• It is hard to learn PL knowledge by end-to-end training.



Huge code corpus







C grammar, type system, semantics







# Can we guide neural network to learn PL knowledge?

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

 Guiding NN to learn grammar rule definitions

#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

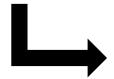
 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

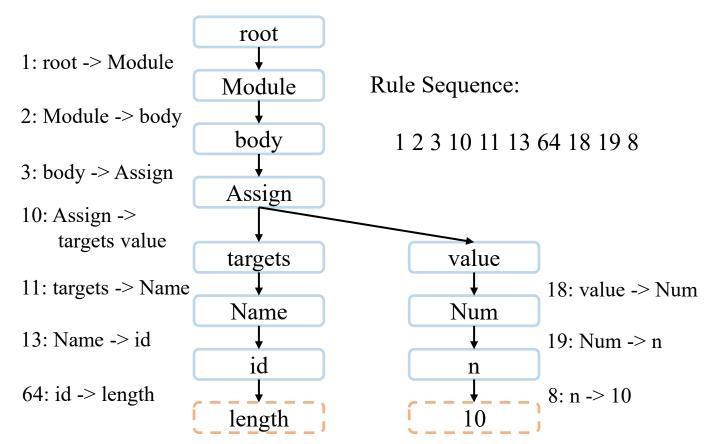
#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## Representing programs as grammar rule lists [TOSEM22]



 Traverse the AST and record the rule ID used to expand each node



## Benefits: Overcoming the problems



- Generate only syntactically correct program
  - if (dp[j][i] == -1) { dp[j][i] = newval; }}}}}}



Much smaller space for more effective learning

Functions between token sequences



Functions between syntactically correct programs

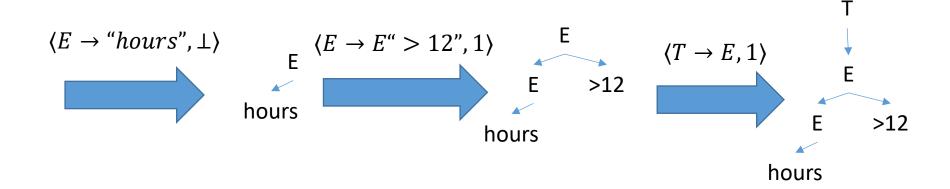


## Benefits:

## COUNTY PROPERTY OF THE PROPERT

## Controlling the generation Order

- Different generation order may lead to significantly different performance
- Generating order can be controlled
  - by extending context-free grammar into expansion grammar



## Benefits: Pruning partial solutions



- Design static program analysis for pruning infeasible partial solutions
- Analyzing partial programs generated by LLM is hard
  - BPE tokens are not even scannable
  - Not to mention parsing
- Analyzing partial AST is much easier
  - if (BoolExpr) then x else x
  - 1 + "x" + Expr

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

 Guiding NN to learn grammar rule definitions

#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## Using Transformer to implement L2S [AAAI20]



- The earliest work that applies Transformer for code generation
  - TreeGen: a Transformer model designed for grammar rule sequences

LPN (Ling et al. 2016) SEQ2TREE (Dong and Lapata 2016) YN17 (Yin and Neubig 2017)	6.1 1.5 16.2	- ~18.2	67.1 53.4
YN17 (Yin and Neubig 2017)		19.2	
` '	16.2	10.2	
1 037 (5 11 11 0 11 0 1771 1 0015)		$\sim 18.2$	75.8
ASN (Rabinovich, Stern, and Klein 2017)	18.2	_	77.6
ReCode (Hayati et al. 2018)	19.6	_	78.4
TreeGen-A	25.8	25.8	79.3
ASN+SUPATT (Rabinovich, Stern, and Klein 2017)	22.7	_	79.2
SZM19 (Sun et al. 2019)	27.3	30.3	79.6
TreeGen-B	31.8	33.3	80.8
	ReCode (Hayati et al. 2018)  TreeGen-A  ASN+SUPATT (Rabinovich, Stern, and Klein 2017)  SZM19 (Sun et al. 2019)	ReCode (Hayati et al. 2018)       19.6         TreeGen-A       25.8         ASN+SUPATT (Rabinovich, Stern, and Klein 2017)       22.7         SZM19 (Sun et al. 2019)       27.3	ReCode (Hayati et al. 2018)       19.6       -         TreeGen-A       25.8       25.8         ASN+SUPATT (Rabinovich, Stern, and Klein 2017)       22.7       -         SZM19 (Sun et al. 2019)       27.3       30.3

TreeGen has been widely applied to decompilation, program repair, code search, automating editing by different researchers

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

 Guiding NN to learn grammar rule definitions

#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

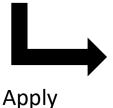
 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

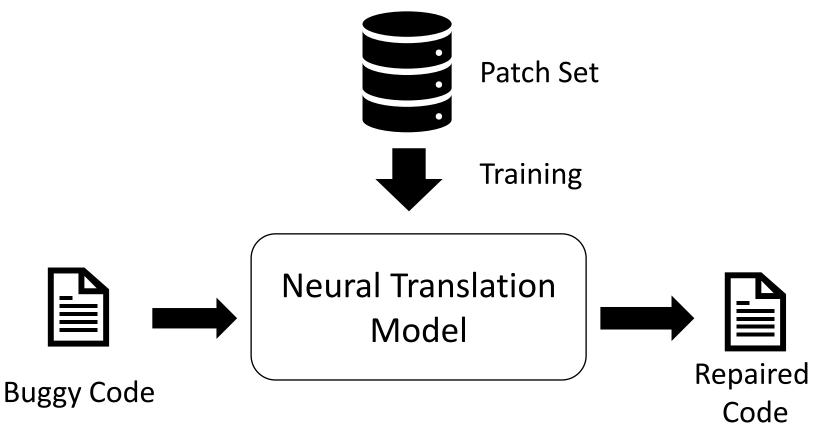
#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## Existing Neural Program Repair



Treating a patch as a pair of code



## A finding in bidirectional transformation [Models'11 MIP]



State-based representation is ineffective

cfa.createEdge(fromNode, Branch.UNCOND, finallyNode);

1. Need to learn diff during training
2. Repr is long (13 tokens)

Delta-based representation is more desirable

modify(9, ON\_EX)

1. Change is directly given 2. Repr is short (3 tokens)

## A grammar of change



```
1. Edits
                           \rightarrow Edit; Edits | end
2. Edit
                           \rightarrow Insert | Modify
3. Insert
                           \rightarrow insert(\langle HLStatement \rangle)
4. Modify
                                  modify(
                                   \langle ID \text{ of an AST Node with a NTS} \rangle,
                                   \langle the same NTS as the above NTS \rangle \rangle
5. \langle Any NTS in HL \rangle \rightarrow
                     copy(\langle ID \ of \ an \ AST \ Node \ with \ the \ same \ NTS \rangle)
                                  ⟨The original production rules in HL⟩
6. (HLIdentifier)
                                  placeholder
                           \rightarrow
                                  (Identifiers in the training set)
```

Ensuring the changed code is still syntactically correct.

## Recoder [ESEC/FSE'21]



- TreeGen for generating changes
- Neural program repair outperformed traditional approaches for the first time

**Table 2: Comparison without Perfect Fault Localization** 

Project	jGenProg	HDRepair	Nopol	CapGen	SketchFix	FixMiner	SimFix	TBar	DLFix	PraPR	AVATAR	Recoder
Chart	0/7	0/2	1/6	4/4	6/8	5/8	4/8	9/14	5/12	4/14	5/12	8/14
Closure	0/0	0/7	0/0	0/0	3/5	5/5	6/8	8/12	6/10	12/62	8/12	17/31
Lang	0/0	2/6	3/7	5/5	3/4	2/3	9/13	5/14	5/12	3/19	5/11	9/15
Math	5/18	4/7	1/21	12/16	7/8	12/14	14/26	18/36	12/28	6/40	6/13	15/30
Time	0/2	0/1	0/1	0/0	0/1	1/1	1/1	1/3	1/2	0/7	1/3	2/2
Mockito	0/0	0/0	0/0	0/0	0/0	0/0	0/0	1/2	1/1	1/6	2/2	2/2
Total	5/27	6/23	5/35	21/25	19/26	25/31	34/56	42/81	30/65	26/148	27/53	53/94
P(%)	18.5	26.1	14.3	84.0	73.1	80.6	60.7	51.9	46.2	17.6	50.9	56.4

In the cells, x/y:x denotes the number of correct patches, and y denotes the number of patches that can pass all the test cases.

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

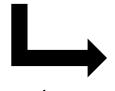
 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## LEAM [ASE'22 Distinguished]



- From Junjie Chen and Lingming Zhang's group
- Exchange the input and output of Recoder
- Program Repairer -> Bug Seeder

Table 4: Overall effectiveness in mutation-based FL

FL	Tech.	Top-1	Top-3	Top-5	MFR	MAR
	Major	35	92	114	9.56	12.42
Metallaxis	PIT	56	102	128	8.16	11.83
Metallaxis	DM	19	47	98	16.64	20.65
	LEAM	118	182	188	3.86	4.57
	Major	35	89	111	10.99	13.11
MUSE	PIT	52	97	124	9.15	11.72
	DM	18	53	94	18.70	22.47
	LEAM	126	181	189	3.88	5.05

## Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

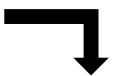
#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## Limit of L2S



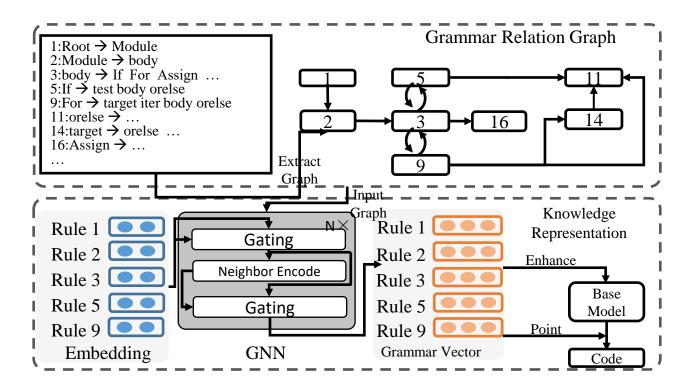
- Force syntactical and other constraint from outside
- NN does not learn their definitions

```
ifstmt -> 'if' '(' boolExpr ')' stmt 10
whilestat -> 'while' '(' boolExpr ')' stmt 11
boolExpr -> andExpr 12
boolExpr -> orExpr 13
```

Grammar rules are encoded as numbers without content. NN could predict impossible sequences such as 10, 11.

## Learning Grammar Rules [IJCAI22]

- Guide the NN to learn grammar definitions
- Word2Vec: assign each token a vector
- Grape: assign each grammar rule a vector, learned with its definition structure



## Learning Grammar Rules [IJCAI22]

- Improve the performance of TreeGen up to 5 percentage points
- Outperforms larger pre-training models

			Semanti	Semantic Parsing						
	Method	HearthStone			Django	Concode	Atis	Job	StrReg	
	Metric	StrAcc	BLEU	Acc+	StrAcc	StrAcc	ExeAcc	ExeAcc	DFAAcc	
	KCAZ13 [Kwiatkowski et al., 2013]	-	-	_	-	-	89.0	-	-	
	WKZ14 [Wang et al., 2014]	-	-	-	-	-	91.3	90.7	-	
ks	SEQ2TREE [Dong and Lapata, 2016]	-	-	_	_	-	84.6	90.0	-	
Networks	ASN+SUPATT [Rabinovich et al., 2017]	22.7	79.2	-	-	-	85.9	92.9	-	
etv	TRANX [Yin and Neubig, 2018]	-	-	-	73.7	-	86.3	90.0	_	
Z	Iyer-Simp+200 idoms [Iyer et al., 2018]	-	-	-	-	12.20	-	-	-	
ıra	GNN-Edge [Shaw et al., 2019]	-	-	-	-	-	87.1	-	-	
Neural	SoftReGex [Park et al., 2019]	-	-	-	_	-	-	-	28.2	
~	TreeGen [Sun et al., 2020]	30.3±1.061	80.8	33.3	76.4	16.6	89.6±0.329	91.5±0.586	22.5	
	GPT-2 [Radford et al., 2019]	16.7	71	18.2	62.3	17.3	84.4	92.1	24.6	
	CodeGPT [Lu et al., 2021]	27.3	75.4	30.3	68.9	18.3	87.5	92.1	22.49	
	TreeGen + Grape	33.6±1.255	85.4	36.3	77.3	17.6	92.16±0.167	92.55±0.817	28.9	

Parameters: TreeGen+Grape: 35M GPT-2、CodeGPT: 110M

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

 Guiding NN to learn grammar rule definitions

#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

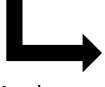
 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## Learning Typing Rules [ICSE23]



Full type system is difficult to learn from data

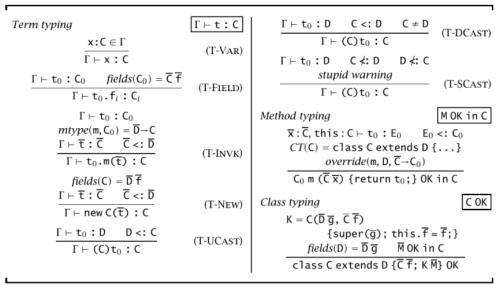


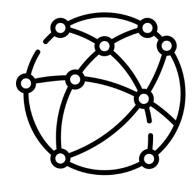
Figure 19-4: Featherweight Java (typing)

Only 30%-40% programs generated by Recoder is typable

## Learning Typing Rules [ICSE23]



- A single rule is much easier to learn
  - T-Graph: present the input of a typing rule to the NN
  - T-Grammar: force NN to predict the output of a typing rule



T-Graph: Representing typing relations

- types of AST nodes
- types of variables
- subtyping relations



T-Grammar:

E -> E && E becomes [Bool]E -> [Bool]E && [Bool]E

## Learning Typing Rules [ICSE23]



#### Applying to program repair, forming Tare

Project	Bugs	CapGen	SimFix	TBar	DLFix	Hanabi	Recoder	Recoder-F	Recoder-T	Tare
Chart	26	4/4	4/8	9/14	5/12	3/5	8/14	9/15	8/16	11/16
Closure	133	0/0	6/8	8/12	6/10	-/-	13/33	14/36	15/31	15/29
Lang	64	5/5	9/13	5/14	5/12	4/4	9/15	9/15	11/23	13/22
Math	106	12/16	14/26	18/36	12/28	19/22	15/30	16/31	16/40	19/42
Time	26	0/0	1/1	1/3	1/2	2/2	2/2	2/2	2/4	2/4
Mockito	38	0/0	0/0	1/2	1/1	-/-	2/2	2/2	2/2	2/2
Total	393	21/25	34/56	42/81	30/65	28/33	49/96	52/101	54/116	62/115

Tare+ExpressAPR(efficient patch validation tool) got the first place in the Java functional bug track of APR-COMP'24.

### Overview



#### TreeGen [AAAI20]

- Using transformer to implement L2S
- The first transformer-based code generator

#### Grape [IJCAI22]

 Guiding NN to learn grammar rule definitions

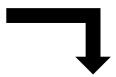
#### Tare [ICSE23]

Guide NN to learn typing rules

#### GrammarT5 [ICSE24]

 Integrating L2S with pretraining`

#### **Implements**



#### L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy pruning



Apply

#### ACS [ICSE17]

 First program repair approach whose precision > 70%

#### Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

#### OCoR [ASE20]

 Code search engine significantly outperforming existing ones

#### LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

## A Era of LLMs









LLMs (=pretrained large models) exhibit superior performance

Can we use grammar-based representation in LLMs?

## Challenges



- Big vocabulary
  - User-defined identifiers can be added to the grammar when the training set is small
  - Pre-training sets are too large
- Heterogeneous grammars
  - Existing models: One programming language
  - Pretraining models: Many programming languages
- Pretraining Tasks
  - Self-supervised training tasks are needed
  - Tasks are expected to guide the neural network to learn the grammar structure

## Big vocabulary



- Existing approaches
  - IDEN -> isodd | iseven
- Our approach
  - Using BPE (Byte Pair Encoding) to find a small set of subtokens
    - is, odd, even
  - Integrating them into the grammar
    - IDEN -> is IDEN | odd IDEN | even IDEN | #is | #odd | #even
    - # indicates the ending tokens
    - Leads to significantly shorter encoding than the standard sequence encoding
      - IDEN -> is IDEN | odd IDEN | even IDEN |  $\epsilon$

## Heterogeneous grammars



- A hyper grammar that includes all grammars
  - Root -> Root@Python | Root@Java | ...
- Experimentally has better performance than sharing some of non-terminals
  - While -> while '(' BoolExpr ')' Statements
  - BoolExpr -> BoolExpr@Java | BoolExpr@C# | ...

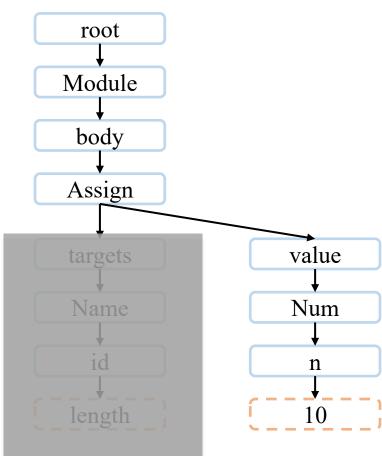
## Pretraining Tasks



Given a rule sequence, predicting the parent of a rule

• 1 2 3 10 11 13 64 18 19 8

 Predicting some subtree of an AST



## Learning Declarations



- Existing pre-training models sort files randomly
- LLMs may see a function or a variable before its declaration
- Dependency parsing:
  - Extract declaration-use relationship from files
  - Sort the files so that declarations appear before use





	Natural-Language-Based Code Generation									
Models	Concode			Conala		Django		MBPP	MathQA	
Metric	BLEU	EM	C-BLEU	BLEU	EM	BLEU	EM	pass@80	pass@80	
TreeGen + Grape(35M)	26.45	17.60	30.05	20.16	2.80	75.86	77.30	2.00	26.58	
GPT-C(110M)	30.85	19.85	33.10	30.32	4.80	72.56	68.91	10.40	58.94	
CodeGPT-adapted(110M)	35.94	20.15	37.27	31.04	4.60	71.24	72.13	12.60	55.90	
CoTexT(220M)	19.19	19.72	38.13	31.45	6.20	75.91	78.43	14.00	58.18	
PLBART(220M)	36.69	18.75	38.52	32.44	5.10	72.81	79.12	12.00	57.25	
CodeT5-small(60M)	38.13	21.55	41.39	31.23	6.00	76.91	81.77	19.20	61.58	
CodeT5-base(220M)	40.73	22.30	43.2	38.91	8.40	81.40	84.04	24.00	71.52	
CodeT5-large(770M)	42.66	22.65	45.08	39.96	7.40	82.11	83.16	32.40	83.14	
Unixcoder(110M)	38.73	22.65	40.86	36.09	10.20	78.42	75.35	22.40	70.16	
GrammarT5-small(60M)	38.68	21.25	41.62	39.18	8.00	81.20	82.77	26.00	84.91	
GrammarT5-base(220M)	42.30	24.75	45.38	41.42	10.40	82.20	84.27	33.20	87.46	

## Industrial Applications

∮ EvalPlus Tests ∮





#	Model	pass@1
1	<b>▼</b> GPT-4-Turbo (April 2024) →	<b>♦</b> 86.6
2	<b>▼</b> <u>DeepSeek-Coder-V2-Instruct</u>	<b>♦</b> 82.3
3	<b>▼</b> GPT-4-Turbo (Nov 2023) *	<b>♦</b> 81.7
4	GPT-4 (May 2023) →	<b>→</b> 79.3
5	CodeQwen1.5-7B-Chat ↔	<b>→</b> 78.7
6	claude-3-opus (Mar 2024).	<b>→</b> 77.4
7	DeepSeek-Coder-33B-instruct ↔	<b>∳</b> 75
8	<u>OpenCodeInterpreter-DS-33B</u> → ♥	<b>∳</b> 73.8
9	WizardCoder-33B-V1.1.  →	<b>→</b> 73.2
10	Artigenz-Coder-DS-6.7B →	<b>4</b> 72.6



One of the best open source LLM for code

Development leaded by my Ph.D. student Qihao Zhu

Dependency parsing was applied

Work-in-progress

Cooperating with a large company to train a LLM with grammar-based representation.

## Industrial Applications





		Total (574 bugs)					
Organization •	Model \$	Plausible @1 ▼	AST Match @1 🕏	Cost \$			
Anthropic	claude-3-5-sonnet-20240620	39,1%	11,7%	\$88.11			
Google	gemini-1.5-pro-002	33,2%	14,3%	\$30.98			
OpenAl	gpt-4o-2024-08-06	31,7%	8,3%	\$30.51			
Google	gemini-1.5-pro-001	28,2%	12,5%	\$78.65			
<b>■</b> Meta	llama-3.1-405b-instruct	27,0%	7,6%	\$29.28			
DeepSeek	deepseek-v2.5	25,1%	6,5%	\$19.73			
Alibaba Cloud	qwen-2.5-72b-instruct	24,2%	6,6%	\$4.74			
Mistral	mistral-large-2407	23,0%	6,6%	\$47.70			
Nvidia	llama-3.1-nemotron-70b-instruct	21,4%	3,4%	\$4.08			
OpenAl	o1-preview-2024-09-12*	N/A	N/A	N/A			



One of the best open source LLM for code

Development leaded by my Ph.D. student Qihao Zhu

Dependency parsing was applied

Work-in-progress

Cooperating with a large company to train a LLM with grammar-based representation.

## Conclusion



- Anxiety: what should we software researchers do if LLMs learn everything by themselves?
- LLMs do not learn programming languages knowledge by themselves
- Guiding them to learn improves their performance
- Future: more genetic ways to learn more software knowledge