

SmartFL: Semantics Based Probabilistic Fault Localization

Yiqian Wu, Zeng Muhan, Yujie Liu, Yi Yin,
Xin Zhang, **Yingfei Xiong**, and Lu Zhang



北京大學
PEKING UNIVERSITY

Status Update 1

- Our group have been working on solving competitive programming problems for the last four years
 - Without using neural networks
 - Talked about this in Meetings 67 and 69
- Our plan: to win a silver medal in the national competition in 2025
- Sep, 2024: o1-IOI won a silver medal in IOI
- Dec, 2024: o3 won a gold medal in IOI, Top 200 in Codeforce
- We surrendered.
- Switched to use LLMs to solve SE problems.
 - Program repair / generation / analysis / verification

Status Update 2

- The grammar-based code representation for neural models
 - which I introduced in
 - IFIP meeting 68
 - Dagstuhl on program synthesis and repair
 - was evaluated in 1.3B and 1.5B models in Kuaishou
 - 4.2-17.7 percentage points higher than token-based representation in accuracy
- The preprint will be released soon

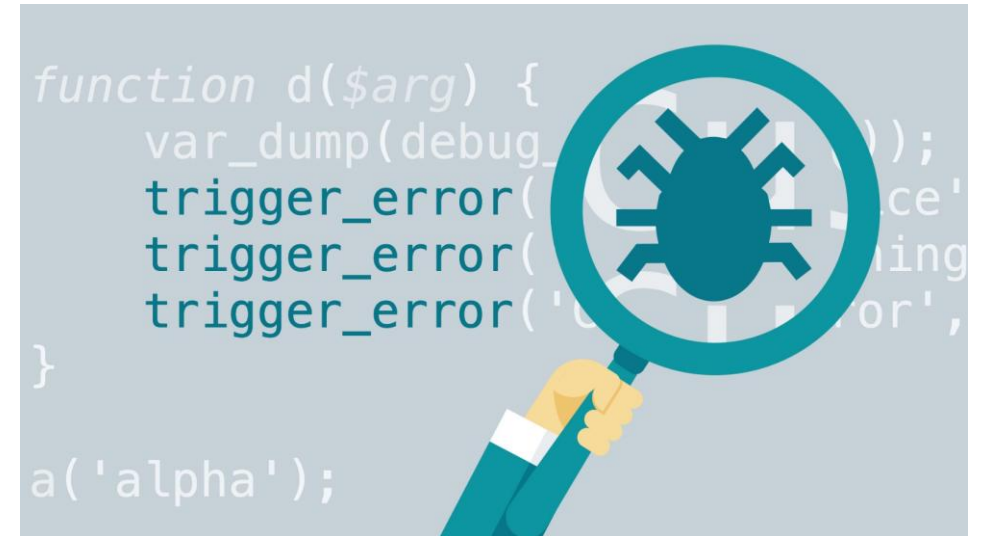
Fault Localization

Input

- A faulty program
- A set of tests with at least one failing test

Output

- The suspiciousness score of each program element



Coverage-based Fault Localization

Idea

- An element covered more by failing tests rather than passing tests is more likely to be faulty

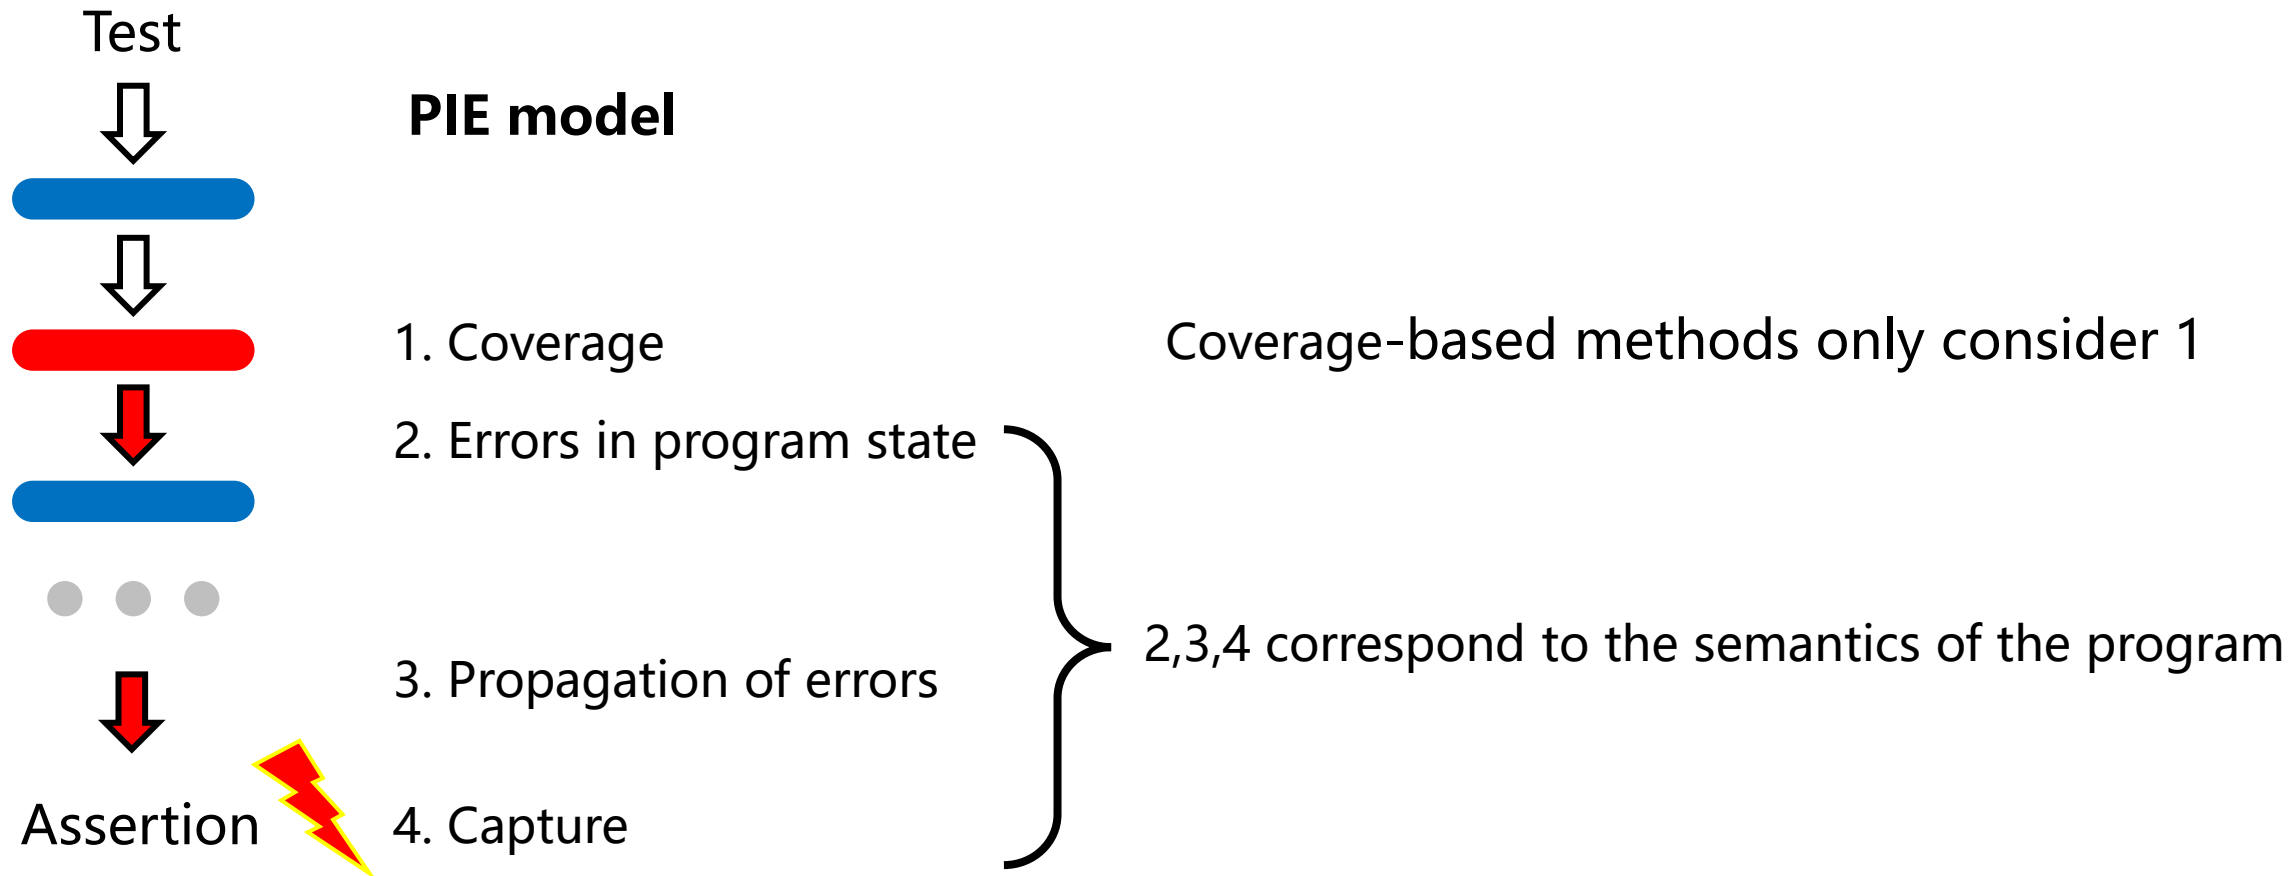
Spectrum-based fault localization

- Count the number of passing/failing tests covering the element
- Calculate the suspiciousness score of each program element

```
foo ( ) {  
    int x, y, z, m, ret ;  
s1  read (x, y, z);  
s2  x = x/y; // correct: x = x * y  
s3  m = x + y;  
s4  if (x > 1) {  
s5      m = x - 2;  
s6      x = x * y;  
s7      z = 2 * y;  
    }  
s8  if (m > 0) {  
s9      ret = x - z;  
    } else {  
s10     ret = y - z;  
    }  
s11 return ret;  
}
```

coverage of 8 tests							
t ₁	t ₂	t ₃	t ₄	t ₅	t ₆	t ₇	t ₈
•	•	•	•	•	•	•	•
•	•	•	•	•	•	•	•
•	•	•	•	•	•	•	•
•	•	•	•	•	•	•	•
	•	•	•	•	•	•	•
	•	•	•			•	•
	•	•	•			•	•
•	•	•	•	•	•	•	•
•		•		•		•	•
	•		•		•		
•	•	•	•	•	•	•	•
results							
F	F	P	P	P	P	F	P

How a Buggy Element Causes the Failure



How to model the latter three conditions?



Existing Research

Mutation-based fault localization

- Generates many mutations on each element
- Watches whether the test result changes
- If a change is more likely to change the results in failing tests, and less likely in passing tests, the corresponding statement is likely to be faulty

Angelic debugging (by Satish et al.)

- Uses symbolic analysis
- Modify the result of an expression
- If an expression can be modified to reverse the results of failing tests while maintaining the results of passing tests, the expression is considered more likely to be faulty

Full Program Semantics  **Heavy**

Our Approach: SmartFL

How a fault leads to the current test results

- The probability of each statement introducing an error
- The probability of each statement propagating an error

Abstracts the full program semantics

- Introduce random variables to represent the correctness of each statement and runtime variable value
- Transform program semantics into probabilistic constraints between random variables

Efficient Probabilistic Modeling of Program Semantics



Running Example

```
1 public class CondTest {  
2     public static int foo(int a) {  
3         if (a <= 2) { // buggy, should be a < 2  
4             a = a + 1;  
5         }  
6         return a;  
7     }  
8  
9     @Test  
10    void pass() {  
11        assertEquals(2, foo(1));  
12    }  
13  
14    @Test  
15    void fail() {  
16        assertEquals(2, foo(2));  
17    }  
18 }
```

$$P(S_3 = 1) = 0.5 \quad P(S_4 = 1) = 0.5 \quad P(V_{p,2} = 1) = 1$$

$$P(V_{p,3} = 1 \mid S_3 = 1 \wedge V_{p,2} = 1) = 1$$

$$P(V_{p,3} = 1 \mid S_3 = 0 \vee V_{p,2} = 0) = 0.5$$

$$P(V_{p,4} = 1 \mid S_4 = 1 \wedge V_{p,2} = 1 \wedge V_{p,3} = 1) = 1$$

$$P(V_{p,4} = 1 \mid S_4 = 0 \vee V_{p,2} = 0 \vee V_{p,3} = 0) = 0.01$$

$$P(S_3 = 0 \mid V_{p,4} = 1 \wedge V_{4,f} = 0) \approx 0.707$$

$$P(S_4 = 0 \mid V_{p,4} = 1 \wedge V_{4,f} = 0) \approx 0.270$$

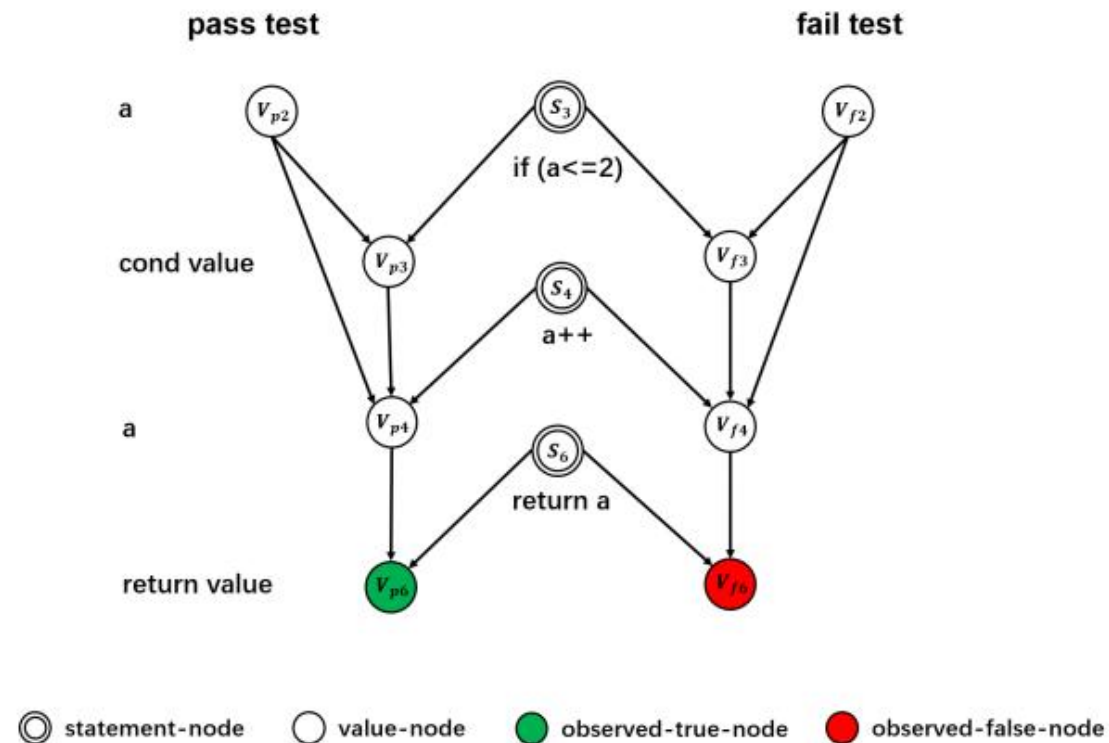
Bernoulli random variables:

S_i : representing the correctness of line i

$V_{t,i}$: representing whether the value defined at line i is correct for test $t \in \{p, f\}$

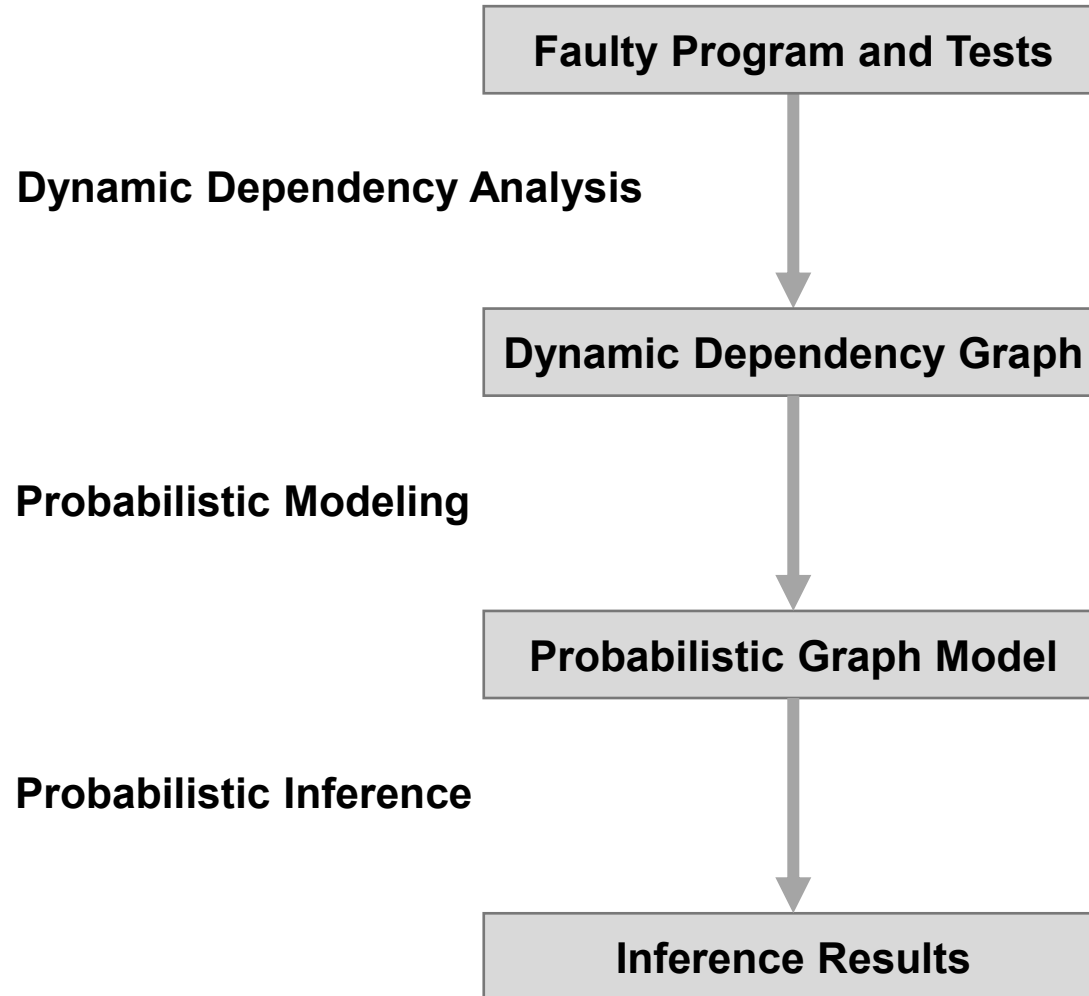
Bayesian Network

```
1 public class CondTest {  
2     public static int foo(int a) {  
3         if (a <= 2) { // buggy, should be a < 2  
4             a = a + 1;  
5         }  
6         return a;  
7     }  
8  
9     @Test  
10    void pass() {  
11        assertEquals(2, foo(1));  
12    }  
13  
14    @Test  
15    void fail() {  
16        assertEquals(2, foo(2));  
17    }  
18 }
```



SmartFL Workflow [ICSE22]

Steps



Results [ICSE22]

- Results on 4 projects from Defects4j 1.0
 - Outperforming both SBFL and MBFL
 - Time consumption between SBFL and MBFL

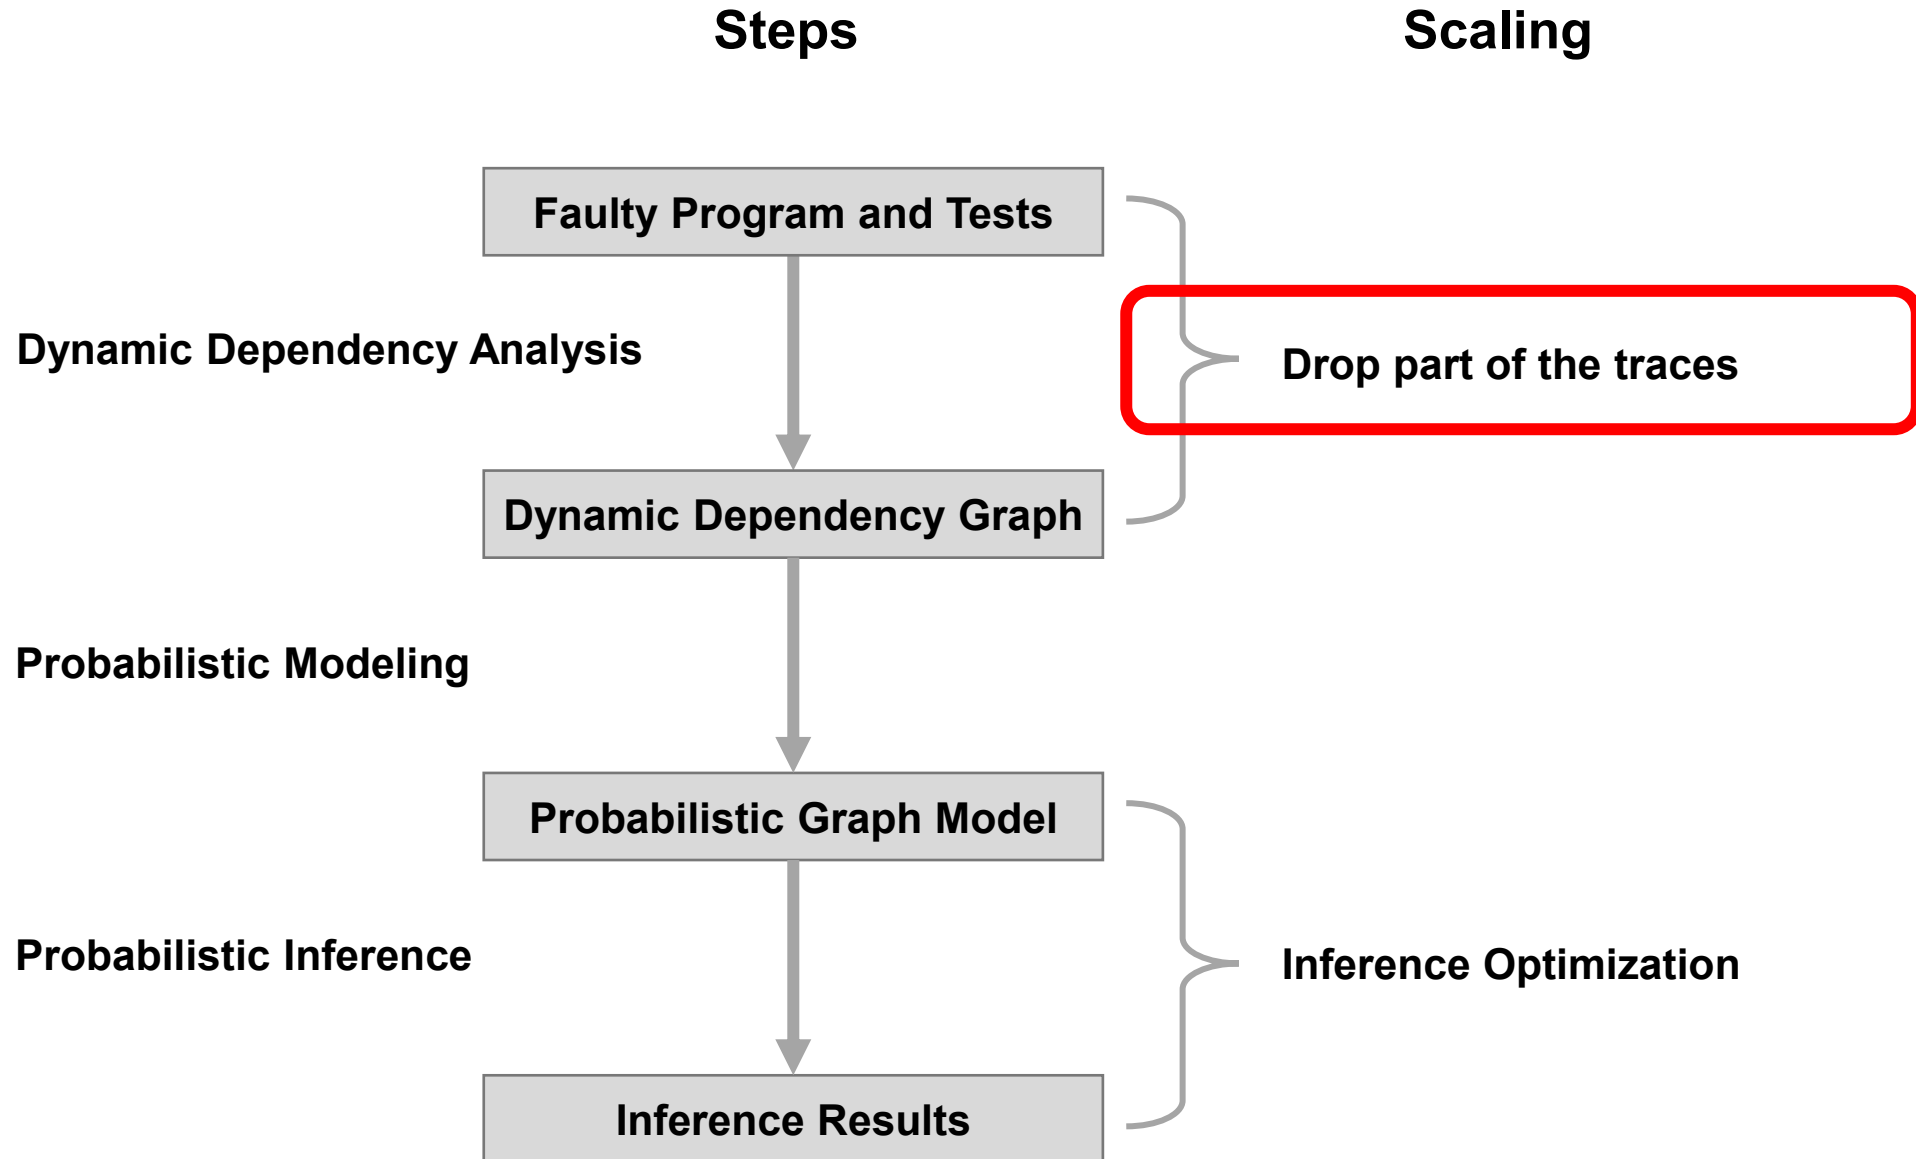
Project	Technique	Top-1	Top-3	Top-5	Top-10
Total	Ochiai	11(5%)	64(29%)	86(39%)	118(53%)
	DStar	12(5%)	65(29%)	86(39%)	117(53%)
	Metallaxis	21(9%)	69(31%)	89(40%)	111(50%)
	MUSE	17(8%)	35(15%)	45(20%)	50(23%)
	SmartFL	47(21%)	80(36%)	97(44%)	118(53%)

Technique	Average	Lang	Math	Chart	Time
Ochiai	64	26	86	44	85
DStar	64	26	86	44	85
Metallaxis	3500	270	3000	5400	12000
MUSE	3500	270	3000	5400	12000
SmartFL	210	51	140	280	830

Bottleneck on Performance [In Submission]


- Traces are too long
- Probabilistic models would be too large for existing probabilistic
 - Drop part of the traces
 - Optimize the probabilistic inference algorithm

SmartFL Workflow [In Submission]




Drop part of traces


Methods not covered by the failing test



Drop large passing tests




Drop large methods in failing tests



Reducing Redundant Methods


Atomic statement

S: a = untraced(b);



Side effects

S1: List l = new ArrayList();
S2: l.add(1); // untraced lib method
S3: x = l.size(); // untraced lib method





Assuming access to and only to the object referenced

Virtual call edge

S1: a = untracedlib(b);
overloaded "equals()" of b

● ● ●
"equals()" return



Reducing Redundant Loops

Same graph structure

Loop



```
S1: while(i<n){
S2:   s = s+i;
S3:   i = i +1;
}
```

Sequence

$(S1\ S2\ S3)^n\ S1$



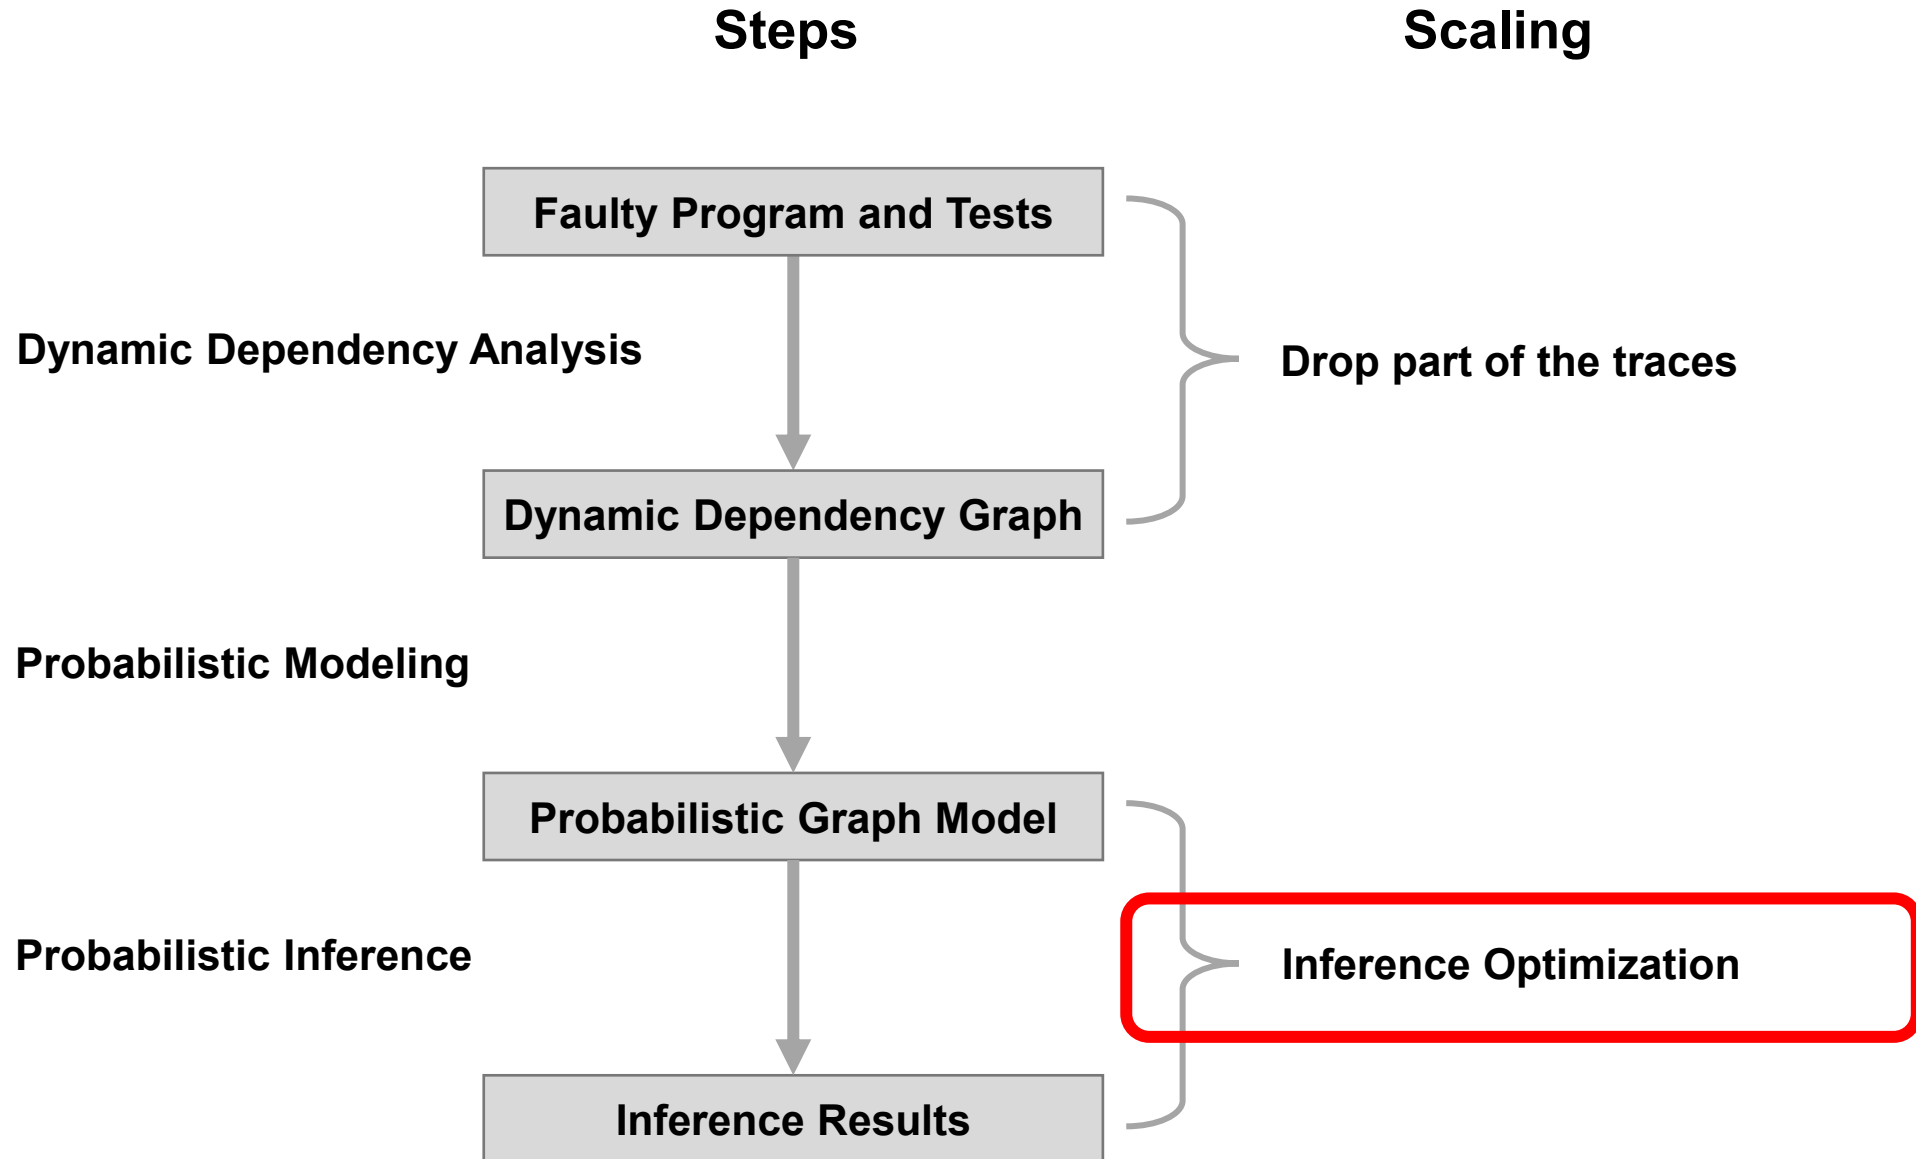
$(S1\ S2\ S3)\ S1$



...

$(S1\ S2)^{100}\ S1S3\ (S1\ S2)^{100} \Rightarrow (S1\ S2)\ S1S3\ (S1\ S2)$

SmartFL Workflow



Inference Optimization

Tabular encoding

$$p(x_v = \text{true} \mid x_1, x_2, \dots, x_n) = \begin{cases} 1, & x_1 \wedge x_2 \cdots \wedge x_n = \text{true} \\ p_0, & x_1 \wedge x_2 \cdots \wedge x_n = \text{false} \end{cases}$$

complexity = $O(2^n)$

Local structure

$$\mathbf{X} = x_1 \wedge x_2 \cdots \wedge x_n \qquad p(x_v = \text{true} \mid \mathbf{X}) = \begin{cases} 1, & \mathbf{X} = \text{true} \\ p_0, & \mathbf{X} = \text{false} \end{cases}$$

complexity = $O(n)$

Effectiveness of SmartFL

Benchmark

Table I: Projects from Defects4J dataset, version 2.0.0.

Project	Faults	LoC	ATests	CTests
Chart	26	203.0k	1818	38
Cli	39	5.7k	262	42
Closure	174	138.8 k	7027	18
Codec	18	10.9k	440	32
Collections	4	67.0k	15582	34
Compress	47	31.0k	432	40
Csv	16	3.1k	180	39
Gson	18	14.0k	988	33
JacksonCore	26	34.4k	356	41
JacksonDatabind	112	95.8k	1610	13
JacksonXml	6	7.6k	152	40
Jsoup	93	15.0k	454	18
JXPath	22	29.2k	305	12
Lang	64	52.3k	1815	30
Math	106	116.2k	3343	29
Mockito	38	18.8k	1156	5
Time	26	67.7k	3802	21
Total	835	76.3k	2604	24

‘Faults’ denotes the number of defective versions of the project, ‘LoC’ denotes the average lines of code of each project, ‘ATests’ denotes the average test numbers of each project, and ‘CTests’ denotes the average number of chosen tests after reducing redundant tests.

Results

Table II: Statement-level Performance

Technique	Top-1	Top-3	Top-5	Top-10
Ochiai	50(6%)	142(17%)	182(22%)	254(30%)
DStar	42(5%)	114(14%)	146(17%)	214(26%)
Metallaxis	51(6%)	132(16%)	166(20%)	219(26%)
MUSE	36(4%)	75(9%)	95(11%)	121(14%)
SmartFL	115(14%)	200(24%)	238(29%)	279(33%)

Table VII: Comparing SmartFL with CAN and UNITE

Technique	Top-1	Top-3	Top-5
CAN	$\leq 15(7\%)$	$\leq 64(28\%)$	$\leq 93(41\%)$
UNITE	$\leq 26(12\%)$	$\leq 75(33\%)$	$\leq 100(45\%)$
SmartFL	47(21%)	88(39%)	103(46%)

Efficiency of SmartFL

Table VIII: Average Time Consumption of each Technique (in seconds)

SBFL	MBFL	SmartFL			
		(a)	(b)	(c)	total
413	46749	41	126	37	205

- a) Profiling (coarse-grained instrumentation to get method-level coverage)
- b) Tracing (getting fine-grained traces of selected tests)
- c) Modeling (building the probabilistic graph and probabilistic inference)

On combining with LLMs

- Neural networks utilize informal information sources, e.g., method names
- SmartFL better captures formal semantic connections
- Potential to be combined
- Attempts
 - Use LLM scores as prior probabilities in SmartFL
 - LLM scores are not really probabilities and tend to dominate
 - Use SmartFL scores in LLMs
 - LLMs do not know how to use them

Table IV: Method-level Performance

Technique	Top-1	Top-3	Top-5	Top-10
Ochiai	167(20%)	305(37%)	351(42%)	398(48%)
DStar	157(19%)	274(33%)	316(38%)	371(44%)
Metallaxis	143(17%)	261(31%)	301(36%)	351(42%)
MUSE	90(11%)	158(19%)	188(23%)	220(26%)
GRACE	280(34%)	382(46%)	438(52%)	\
SmartFL	213(26%)	326(39%)	372(45%)	424(51%)

Table V: Comparing SmartFL with LEAM

Technique	Top-1	Top-3	Top-5
LEAM-Metallaxis	118(53%)	182(81%)	188(84%)
LEAM-MUSE	126(56%)	181(81%)	189(84%)
SmartFL	91(41%)	131(58%)	149(67%)

Conclusion

Main Contributions

1. A fault localization approach by efficient approximation of program semantics.
2. Novel techniques to reduce the size of the model and to efficiently infer posterior probabilities for addressing the scalability challenge.
3. An evaluation on the Defects4J dataset to show the effectiveness and the efficiency of our approach.

Tool and Data

<https://github.com/toledosakasa/SMARTFL>