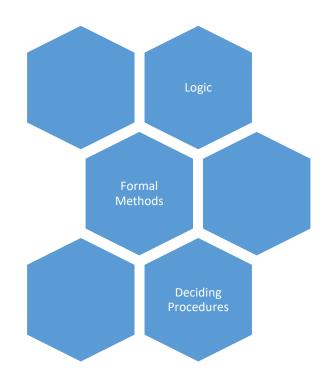


# Neural Code Generation Models with Programming Language Knowledge

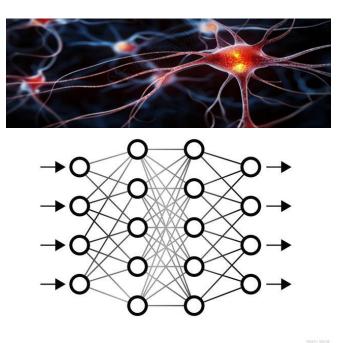
Yingfei Xiong Peking University

# Two Approaches to Al





Symbolism
Learn from humans

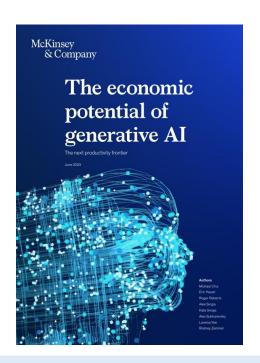


Connectionism

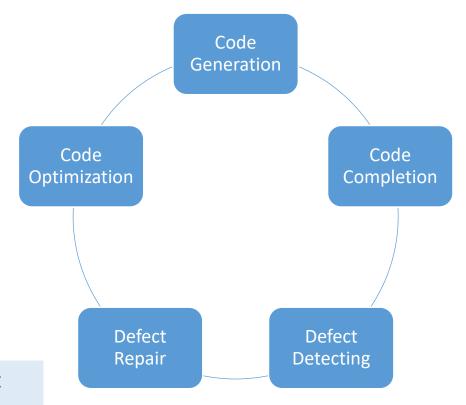
Learn from human creators

# Important LLM Application Assisting Software Development





McKinsey: save 20%-45% development cost











# Mismatch: connectionism



# and programs

Programs are symbolic product, with symbolic knowledge:

- Syntactic: ()+5(illegal
- Type: 1+true illegal
- Semantic: use without initialization illegal

These symbolic knowledges are difficult to be learned by neural network

# How Secure is Code Generated by ChatGPT?

Raphaël Khoury<sup>1</sup>, Anderson R. Avila<sup>2</sup>, Jacob Brunelle<sup>1</sup>, Baba Mamadou Camara<sup>1</sup>

<sup>1</sup>Université du Quebec en Outaouais, Quebec, Canada <sup>2</sup>Institut national de la recherche scientifique, Quebec, Canada {raphael.khoury, anderson.raymundoavila, bruj30, camb12}@uqo.ca

Abstract-In recent years, large language models have been responsible for great advances in the field of artificial intelligence (AI). ChatGPT in particular, an AI chatbot developed and recently released by OpenAI, has taken the field to the next level. The conversational model is able not only to process human-like text, but also to translate natural language into code. However, the safety of programs generated by ChatGPT should not be overlooked. In this paper, we perform an experiment to address this issue. Specifically, we ask ChatGPT to generate a number of program and evaluate the security of the resulting source code. We further investigate whether ChatGPT can be prodded to improve the security by appropriate prompts, and discuss the ethical aspects of using AI to generate code. Results suggest that ChatGPT is aware of potential vulnerabilities, but nonetheless often generates source code that are not robust to certain attacks. Index Terms-Large language models, ChatGPT, code security,

#### I. INTRODUCTION

automatic code generation

For years, large language models (LLM) have been demonstrating impressive performance on a number of natural language processing (NLP) tasks, such as sentiment analysis, natural language understanding (NLU), machine translation (MT) to name a few. This has been possible specially by means of increasing the model size, the training data and the model complexity [1]. In 2020, for instance, OpenAI announced GPT-3 [2], a new LLM with 175B parameters, 100 times larger than GPT-2 [3]. Two years later, ChatGPT [4], an artificial intelligence (AI) chatbot capable of understanding and generating human-like text, was released. The conversational AI model, empowered in its core by an LLM based on the Transformer architecture, has received great attention from both industry and academia, given its potential to be applied in different downstream tasks (e.g., medical reports [5], code generation 6, educational tool [7], etc).

Therefore, this paper is an attempt to answer the question of how secure is the source code generated by ChatGPT. Moreover, we investigate and propose follow-up questions that can guide ChatGPT to assess and regenerate more secure source code.

In this paper, we perform an experiment to evaluate the security of code generated by ChatGPT, fine-tuned from a model in the GPT-3.5 series. Specifically, we asked ChatGPT to generate 21 programs, in 5 different programming languages: C, C++, Python, html and Java. We then evaluated the generated program and questioned ChatGPT about any vulnerability present in the code. The results were worrisome. We found that, in several cases, the code generated by ChatGPT fell well below minimal security standards applicable in most contexts. In fact, when prodded to whether or not the produced code was secure, ChatGTP was able to recognize that it was not. The chatbot, however, was able to provide a more secure version of the code in many cases if explicitly asked to do so.

The remainder of this paper is organized as follows. Section III describes our methodology as well as provides an overview of the dataset. Section IIII details the security flaws we found in each program. In Section IVI we discuss our results, as well as the ethical consideration of using AI models to generate code. Section IVI surveys related works. Section IVI discusses threats to the validity of our results. Concluding remarks are given in Section IVII.

#### II. STUDY SETUP

#### A. Methodology

In this study, we asked ChatGPT to generate 21 programs, using a variety of programming languages. The programs generated serve a diversity of purpose, and each program was chosen to highlight risks of a specific vulnerability (eg.

76% of the programs generated by GPT contain vulnerabilities

# Symbolic PL knowledge is useful



Correct code generation requires symbolic PL

knowledges

bool and(bool a, bool b) {
 \_\_\_\_\_\_
}

assignment is common and should be used

NN that does not know types

all parameters are Boolean so 'if' is more likely

NN that knows types



# Can we guide neural network to learn PL knowledge?

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

### Tare [ICSE23]

•Guide NN to learn typing rules

## GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

# GrammarCoder [ACL25-Finding]

- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis

# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

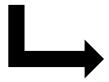
• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

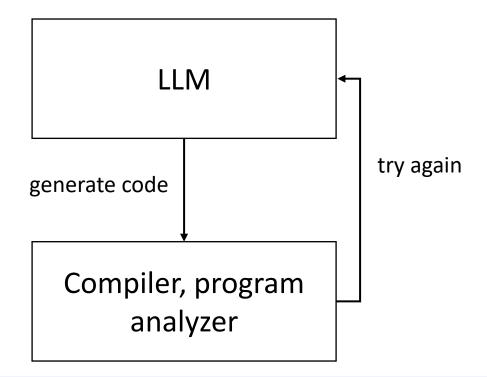


Apply

# Generating only safe code



Attempt 1: Check after generation

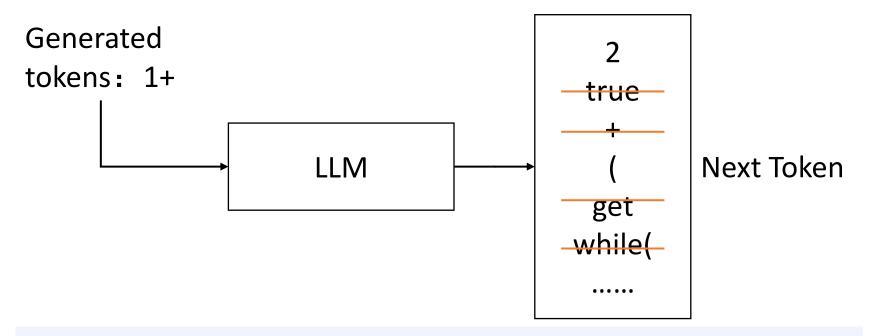


Inefficient, may keep generating code with the same fault

# Generating only safe code



Attempt 2: constrained decoding



**Difficulty:** Generated tokens are defined by the BPE algorithm. Lexical analysis is already difficult, let alone parsing.

# L2S: representing by grammar rule sequences [TOSEM22]



Program

x+y

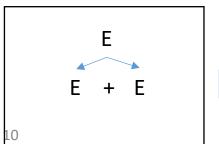
**Token Sequence** 

x, +, y

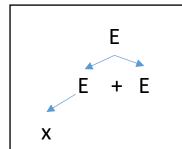
Grammar Rule Sequence

 $r_1, r_2, r_3$ 

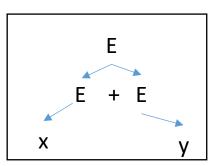
$$r_1: E \to E + E$$
  
 $r_2: E \to x$   
 $r_3: E \to y$ 











# Benefit: Constrained Decoding Made Easy



- Ensuring syntactic correctness is trivial
- Type and semantic analysis can be performed on partial AST
  - 1 + "x" + Expr //Type incorrect
  - if (BoolExpr) then x else x //Semantically incorrect for x=1 → ret=2

# • Step:

- Pre-analysis on grammar rules: get all possibilities for a non-terminal
- Abstract interpretation on partial program

# Benefit: Better Alignment with Semantics [ACL25-Finding]



# Same semantics

```
if (x<0) y=y+1;</li>
if (x < 0) {
    y = y + 1;
}</li>
```

Similar in Grammar Rule Representation Different in Token Representation

# Different semantics

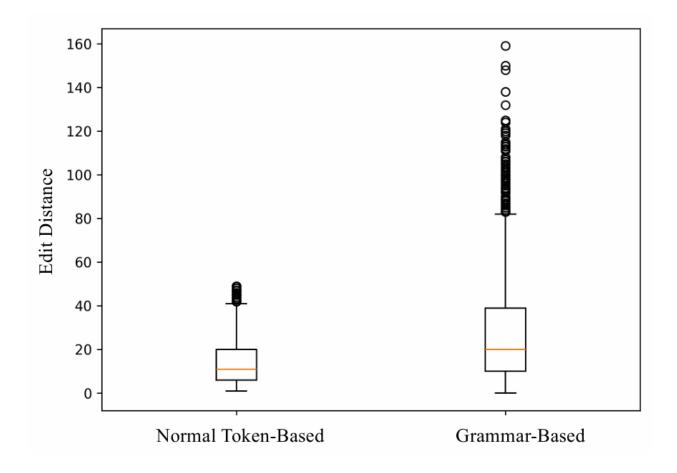
```
    for i in range(1, 6):
    x = x + 1
    sum = sum + x
```

```
    for i in range(1, 6):
    x = x + 1
    sum = sum + x
```

Different in Grammar Rule Representation Similar in Token Representation

# Edit Distance between Semantically Different Code

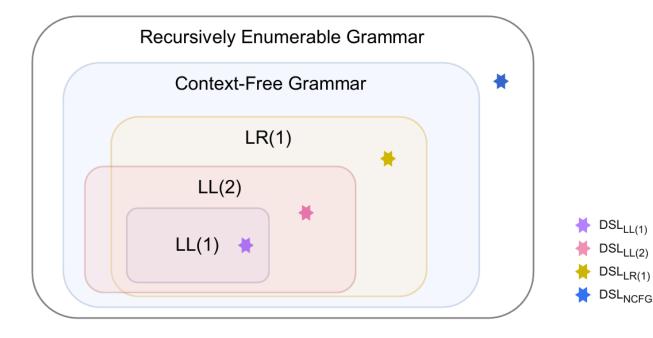


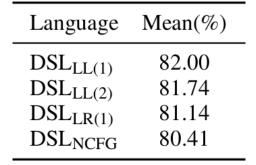


# Benefit: Easier Parsing [Submitted]



 The easier the language is to parse, the better the performance of the neural model.





# Benefit: Easier Parsing [Submitted]



 The easier the language is to parse, the better the performance of the neural model.

If n is an integer and  $101 * n^2 \le 3600$ , what is the greatest possible value of n? n0 = 101.0 n1 = 2.0 n2 = 3600.0

a. Question

- 1. import math 2. n0 = 101.0
- 3. n1 = 2.0
- 4. n2 = 3600.0
- 5. t0 = n2 / n0
- 6. t1 = math.sqrt(max(0, t0))
- 7. answer = math.floor(t1)
  - b. Answer in  $DSL_{LR(1)}$

- 1. import math
- 2. n0 = 101.0
- 3. n1 = 2.0
- 4. n2 = 3600.0
- 5. t0 = / n2 n0
- 6.  $t1 = \langle call \rangle \langle attr \rangle$  math.sqrt( $\langle call \rangle \rangle$  max(0, t0)
- 7. answer = <call> <attr> math.floor(t1)
  - c. Answer in  $DSL_{LL(1)}$

- 1. import math 2. <exp> n0 = 101.0
- 3.  $\langle \exp \rangle$  n1 = 2.0
- 5. \exp> 111 2.0
- 4. <exp> n2 = 3600.0
- 5. < exp > t0 = < exp > / n2 n0
- 6. <exp> t1 = <exp> <call> <exp> <attr> math.sqrt(<exp> <call> max(0, t0))
- 7. <exp> answer = <exp> <call> <exp> <attr> math.floor(t1)
  - d. Answer in DSL<sub>LL(2)</sub>

- 1. import math
- 2. n0 = 101.0 : n0
- 3. n1 = 2.0; n1
- 4. n2 = 3600.0 ; n2
- 5. t0 = n2 / n0 ; t0
- 6. t1 = math.sqrt(max(0, t0)); t1
- 7. answer = math.floor(t1); ans
  - e. Answer in DSL<sub>CSG</sub>

Language	Mean(%)
DSL <sub>LL(1)</sub>	82.00
DSL <sub>LL(2)</sub>	81.74
DSL <sub>LR(1)</sub>	81.14
DSL <sub>NCFG</sub>	80.41

# Benefit: Easier Parsing [Submitted]



- Grammar rule representation is in LL(1)
- Significantly easier than most languages
  - Python: not context-free
  - Java: LR

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

### Tare [ICSE23]

•Guide NN to learn typing rules

### GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

# GrammarCoder [ACL25-Finding]

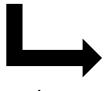
- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis



Apply

# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# Using Transformer to implement L2S [AAAI20]



- The earliest work that applies Transformer for code generation
  - TreeGen: a Transformer model designed for grammar rule sequences

	Model	StrAcc	Acc+	BLEU
iii	LPN (Ling et al. 2016)	6.1	_	67.1
Plain	SEQ2TREE (Dong and Lapata 2016)	1.5	_	53.4
	YN17 (Yin and Neubig 2017)	16.2	$\sim 18.2$	75.8
	ASN (Rabinovich, Stern, and Klein 2017)	18.2	_	77.6
İ	ReCode (Hayati et al. 2018)	19.6	_	78.4
	TreeGen-A	25.8	25.8	79.3
ral	ASN+SUPATT (Rabinovich, Stern, and Klein 2017)	22.7	_	79.2
ctu	SZM19 (Sun et al. 2019)	27.3	30.3	79.6
Structural	TreeGen-B	31.8	33.3	80.8

TreeGen has been widely applied to decompilation, program repair, code search, automating editing by different researchers

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

### Tare [ICSE23]

•Guide NN to learn typing rules

### GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

## GrammarCoder [ACL25-Finding]

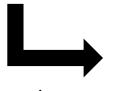
- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis



Apply

# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

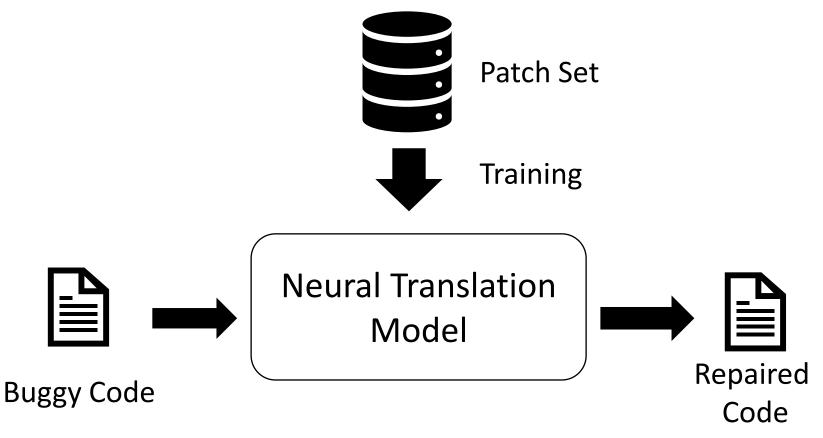
# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# Existing Neural Program Repair



Treating a patch as a pair of code



# A finding in bidirectional transformation [Models'11 MIP]



State-based representation is ineffective

cfa.createEdge(fromNode, Branch.UNCOND, finallyNode);

cfa.createEdge(fromNode, Branch.ON\_EX, finallyNode);

1. Need to learn diff during training
2. Repr is long (13 tokens)

Delta-based representation is more desirable

modify(9, ON\_EX)

1. Change is directly given 2. Repr is short (3 tokens)

# A grammar of change



```
1. Edits
                           \rightarrow Edit; Edits | end
2. Edit
                            \rightarrow Insert | Modify
3. Insert
                            \rightarrow insert(\langle HLStatement \rangle)
4. Modify
                                   modify(
                                    \langle ID \text{ of an AST Node with a NTS} \rangle,
                                    \langle the same NTS as the above NTS \rangle \rangle
5. \langle Any NTS in HL \rangle \rightarrow
                      copy(\langle ID \ of \ an \ AST \ Node \ with \ the \ same \ NTS \rangle)
                                   ⟨The original production rules in HL⟩
6. (HLIdentifier)
                                   placeholder
                            \rightarrow
                                   (Identifiers in the training set)
```

Ensuring the changed code is still syntactically correct.

# Recoder [ESEC/FSE'21]



- TreeGen for generating changes
- Neural program repair outperformed traditional approaches for the first time

**Table 2: Comparison without Perfect Fault Localization** 

Project	jGenProg	HDRepair	Nopol	CapGen	SketchFix	FixMiner	SimFix	TBar	DLFix	PraPR	AVATAR	Recoder
Chart	0/7	0/2	1/6	4/4	6/8	5/8	4/8	9/14	5/12	4/14	5/12	8/14
Closure	0/0	0/7	0/0	0/0	3/5	5/5	6/8	8/12	6/10	12/62	8/12	17/31
Lang	0/0	2/6	3/7	5/5	3/4	2/3	9/13	5/14	5/12	3/19	5/11	9/15
Math	5/18	4/7	1/21	12/16	7/8	12/14	14/26	18/36	12/28	6/40	6/13	15/30
Time	0/2	0/1	0/1	0/0	0/1	1/1	1/1	1/3	1/2	0/7	1/3	2/2
Mockito	0/0	0/0	0/0	0/0	0/0	0/0	0/0	1/2	1/1	1/6	2/2	2/2
Total	5/27	6/23	5/35	21/25	19/26	25/31	34/56	42/81	30/65	26/148	27/53	53/94
P(%)	18.5	26.1	14.3	84.0	73.1	80.6	60.7	51.9	46.2	17.6	50.9	56.4

In the cells, x/y:x denotes the number of correct patches, and y denotes the number of patches that can pass all the test cases.

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

### Tare [ICSE23]

•Guide NN to learn typing rules

### GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

## GrammarCoder [ACL25-Finding]

- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis

Apply

# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# LEAM [ASE'22 Distinguished]



- From Junjie Chen and Lingming Zhang's group
- Exchange the input and output of Recoder
- Program Repairer -> Bug Seeder

Table 4: Overall effectiveness in mutation-based FL

FL	Tech.	Top-1	Top-3	Top-5	MFR	MAR
	Major	35	92	114	9.56	12.42
Matallarria	PIT	56	102	128	8.16	11.83
Metallaxis	DM	19	47	98	16.64	20.65
	LEAM	118	182	188	3.86	4.57
	Major	35	89	111	10.99	13.11
MUSE	PIT	52	97	124	9.15	11.72
MUSE	DM	18	53	94	18.70	22.47
	LEAM	126	181	189	3.88	5.05

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

# Grape [IJCAI22]

•Guiding NN to learn grammar rule definitions

### Tare [ICSE23]

•Guide NN to learn typing rules

### GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

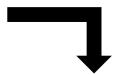
### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

## GrammarCoder [ACL25-Finding]

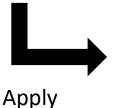
- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis



# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# Limit of L2S



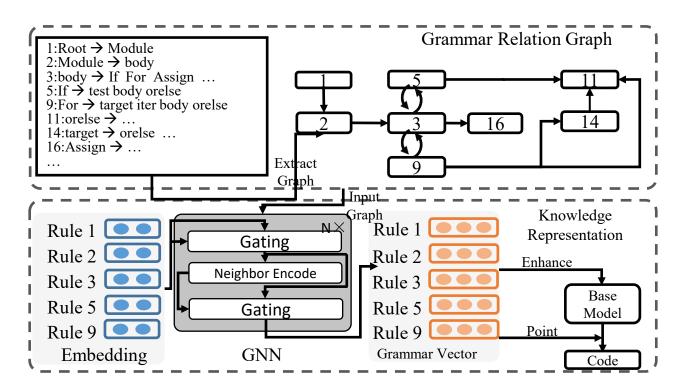
- Force syntactical and other constraint from outside
- NN does not learn their definitions

```
ifstmt -> 'if' '(' boolExpr ')' stmt 10
whilestat -> 'while' '(' boolExpr ')' stmt 11
boolExpr -> andExpr 12
boolExpr -> orExpr 13
```

Grammar rules are encoded as numbers without content. NN could predict impossible sequences such as 10, 11.

# Learning Grammar Rules [IJCAI22]

- Guide the NN to learn grammar definitions
- Word2Vec: assign each token a vector
- Grape: assign each grammar rule a vector, learned with its definition structure



# Learning Grammar Rules [IJCAI22]

- Improve the performance of TreeGen up to 5 percentage points
- Outperforms larger pre-training models

			Code	e Genera	tion	Semanti	Semantic Parsing		
	Method	HearthStone			Django	Concode	Atis	Job	StrReg
	Metric	StrAcc	BLEU	Acc+	StrAcc	StrAcc	ExeAcc	ExeAcc	DFAAcc
	KCAZ13 [Kwiatkowski et al., 2013]	-	-	-	_	-	89.0	-	-
	WKZ14 [Wang et al., 2014]	-	-	-	-	-	91.3	90.7	-
ks	SEQ2TREE [Dong and Lapata, 2016]	-	-	-	-	-	84.6	90.0	-
Neural Networks	ASN+SUPATT [Rabinovich et al., 2017]	22.7	79.2	-	-	-	85.9	92.9	-
etv	TRANX [Yin and Neubig, 2018]	-	-	-	73.7	-	86.3	90.0	-
Z	Iyer-Simp+200 idoms [Iyer et al., 2018]	-	-	-	-	12.20	-	-	-
ıra	GNN-Edge [Shaw et al., 2019]	-	-	-	-	-	87.1	-	-
\e	SoftReGex [Park et al., 2019]	-	-	-	-	-	-	-	28.2
	TreeGen [Sun et al., 2020]	30.3±1.061	80.8	33.3	76.4	16.6	89.6±0.329	91.5±0.586	22.5
	GPT-2 [Radford et al., 2019]	16.7	71	18.2	62.3	17.3	84.4	92.1	24.6
	CodeGPT [Lu et al., 2021]	27.3	75.4	30.3	68.9	18.3	87.5	92.1	22.49
	TreeGen + Grape	33.6±1.255	85.4	36.3	77.3	17.6	92.16±0.167	92.55±0.817	28.9

Parameters: TreeGen+Grape: 35M GPT-2、CodeGPT: 110M

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

# Tare [ICSE23]

•Guide NN to learn typing rules

## GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

### DeepSeek-Coder [arxiv]

- •Guiding NN to learn Def-Use relations
- •SOTA open-source code model at that time

# GrammarCoder [ACL25-Finding]

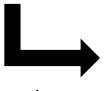
- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis



Apply

# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# Learning Typing Rules [ICSE23]



Full type system is difficult to learn from data

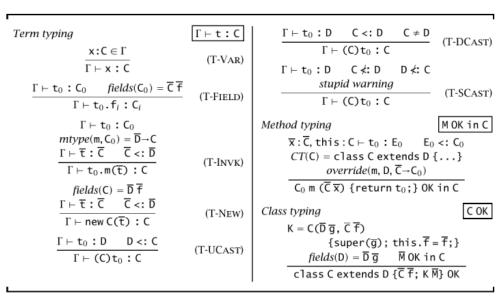


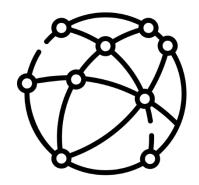
Figure 19-4: Featherweight Java (typing)

Only 30%-40% programs generated by Recoder is typable

# Learning Typing Rules [ICSE23]



- A single rule is much easier to learn
  - T-Graph: present the input of a typing rule to the NN
  - T-Grammar: force NN to predict the output of a typing rule



T-Graph: Representing typing relations

- types of AST nodes
- types of variables
- subtyping relations



T-Grammar: E -> E && E becomes [Bool]E -> [Bool]E && [Bool]E

# Learning Typing Rules [ICSE23]



# Applying to program repair, forming Tare

Project	Bugs	CapGen	SimFix	TBar	DLFix	Hanabi	Recoder	Recoder-F	Recoder-T	Tare
Chart	26	4/4	4/8	9/14	5/12	3/5	8/14	9/15	8/16	11/16
Closure	133	0/0	6/8	8/12	6/10	-/-	13/33	14/36	15/31	15/29
Lang	64	5/5	9/13	5/14	5/12	4/4	9/15	9/15	11/23	13/22
Math	106	12/16	14/26	18/36	12/28	19/22	15/30	16/31	16/40	19/42
Time	26	0/0	1/1	1/3	1/2	2/2	2/2	2/2	2/4	2/4
Mockito	38	0/0	0/0	1/2	1/1	-/-	2/2	2/2	2/2	2/2
Total	393	21/25	34/56	42/81	30/65	28/33	49/96	52/101	54/116	62/115

Tare+ExpressAPR(efficient patch validation tool) got the first place in the Java functional bug track of APR-COMP'24.

# Overview



### TreeGen [AAAI20]

- •Using transformer to implement L2S
- •The first transformer-based code generator
- •SOTA 35m model

### Grape [IJCAI22]

• Guiding NN to learn grammar rule definitions

# Tare [ICSE23]

•Guide NN to learn typing rules

## GrammarT5 [ICSE24]

- Applying L2S to pretraining
- •SOTA 0.2B model at that time

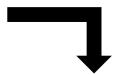
### DeepSeek-Coder [arxiv]

- Guiding NN to learn Declare-Use relations
- •SOTA open-source code model at that time

## GrammarCoder [ACL25-Finding]

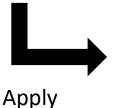
- Applying L2S to Decoder-only
- •SOTA 1.5B code model

# **Implements**



# L2S Framework [TOSEM22]

- Representing code as grammar rule sequences
- Ensuring syntactic correctness
- Allowing easy symbolic analysis



# ACS [ICSE17]

• First program repair approach whose precision > 70%

# Recoder [FSE23]

 First neural program repair approach outperforming traditional approaches

# ET[APRCOMP24]

• 1st Place in APR-COMP

# OCoR [ASE20]

 Code search engine significantly outperforming existing ones

# LEAM [ASE22]

 Mutation generation engine that significantly outperforming existing ones

# A Era of LLMs









LLMs (=pretrained large models) exhibit superior performance

Can we use grammar-based representation in LLMs?

# Challenges



- Big vocabulary
  - User-defined identifiers can be added to the grammar when the training set is small
  - Pre-training sets are too large
- Heterogeneous grammars
  - Existing models: One programming language
  - Pretraining models: Many programming languages
- Pretraining Tasks
  - Self-supervised training tasks are needed
  - Tasks are expected to guide the neural network to learn the grammar structure

# Big vocabulary



- Existing approaches
  - IDEN -> isodd | iseven
- Our approach
  - Using BPE (Byte Pair Encoding) to find a small set of subtokens
    - is, odd, even
  - Integrating them into the grammar
    - IDEN -> is IDEN | odd IDEN | even IDEN | #is | #odd | #even
    - # indicates the ending tokens
    - Leads to significantly shorter encoding than the standard sequence encoding
      - IDEN -> is IDEN | odd IDEN | even IDEN |  $\epsilon$

# Heterogeneous grammars



- A hyper grammar that includes all grammars
  - Root -> Root@Python | Root@Java | ...
- Experimentally has better performance than sharing some of non-terminals
  - While -> while '(' BoolExpr ')' Statements
  - BoolExpr -> BoolExpr@Java | BoolExpr@C# | ...

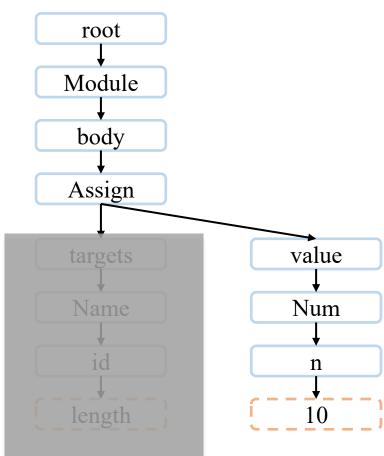
# Pretraining Tasks



Given a rule sequence, predicting the parent of a rule

• 1 2 3 10 11 13 64 18 19 8

 Predicting some subtree of an AST



# Learning Declare-Use Relation



- Existing pre-training models sort files randomly
- LLMs may see a function or a variable before its declaration
- Dependency parsing:
  - Extract declaration-use relationship from files
  - Sort the files so that declarations appear before use





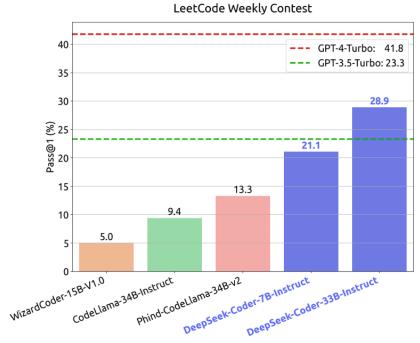
	Natural-Language-Based Code Generation								
Models	Concode			Conala		Django		MBPP	MathQA
Metric	BLEU	EM	C-BLEU	BLEU	EM	BLEU	EM	pass@80	pass@80
TreeGen + Grape(35M)	26.45	17.60	30.05	20.16	2.80	75.86	77.30	2.00	26.58
GPT-C(110M)	30.85	19.85	33.10	30.32	4.80	72.56	68.91	10.40	58.94
CodeGPT-adapted(110M)	35.94	20.15	37.27	31.04	4.60	71.24	72.13	12.60	55.90
CoTexT(220M)	19.19	19.72	38.13	31.45	6.20	75.91	78.43	14.00	58.18
PLBART(220M)	36.69	18.75	38.52	32.44	5.10	72.81	79.12	12.00	57.25
CodeT5-small(60M)	38.13	21.55	41.39	31.23	6.00	76.91	81.77	19.20	61.58
CodeT5-base(220M)	40.73	22.30	43.2	38.91	8.40	81.40	84.04	24.00	71.52
CodeT5-large(770M)	42.66	22.65	45.08	39.96	7.40	82.11	83.16	32.40	83.14
Unixcoder(110M)	38.73	22.65	40.86	36.09	10.20	78.42	75.35	22.40	70.16
GrammarT5-small(60M)	38.68	21.25	41.62	39.18	8.00	81.20	82.77	26.00	84.91
GrammarT5-base(220M)	42.30	24.75	45.38	41.42	10.40	82.20	84.27	33.20	87.46

SOTA 0.2B model in 2024 Standard L2S was applied

# DeepSeek-Coder [arxiv24]







Cooperation with DeepSeek

SOTA open source code model in 2024, still widely used as a base model for fine-tuning

Dependency parsing was applied

# GrammarCoder [ACL25-Finding]

	CUNIDA
. 7	
	.1898.

Model	HumanEval(+)	MBPP(+)
Original		
DeepSeek-Coder-1.3B-Base	34.8 (28.7)	56.7 (47.9)
Qwen2.5-1.5B-Base	37.2 (32.9)	60.2 (49.6)
Normal Token-Based CPT		
DeepSeek-Coder-1.3B-Base (CPT)	43.9 (39.6)	61.4 (51.3)
Qwen2.5-1.5B-Base (CPT)	50.6 (42.7)	60.3 (51.1)
Grammar-Based CPT		
GrammarCoder-1.3B-Base	<b>63.4</b> (57.3)	68.3 (56.9)
GrammarCoder-1.5B-Base	63.4 (59.1)	64.8 (55.3)

Cooperation with Kuaishou SOTA 1.5B model Standard L2S was applied

# GrammarCoder [ACL25-Finding]



Model	HumanEval	HumanEval+	MBPP	MBPP+
DeepSeek-Coder-1.3B-Instruct (Guo et al., 2024)	65.9	60.4	64.3	54.8
Qwen2.5-1.5B-Instruct (Team, 2024)	61.6	49.4	63.2	55.6
OpenCoder-1.5B-Instruct (Huang et al., 2024)	72.5	67.7	72.7	61.9
Yi-Coder-1.5B-Chat (AI et al., 2025)	67.7	63.4	68.0	59.0
Phi-3-Mini-4K-3.8B-Instruct (Abdin et al., 2024)	64.6	59.1	65.9	54.2
CodeGemma-7B-Instruct (Team et al., 2024)	60.4	51.8	70.4	56.9
GrammarCoder-1.3B-Instruct	70.7	64.0	71.2	58.7
GrammarCoder-1.5B-Instruct	73.2	68.3	73.3	61.1

Cooperation with Kuaishou SOTA 1.5B model Standard L2S was applied

# Conclusion



- Code has structure
  - Syntax
  - Types
  - Semantics
- These structures are defined by symbolic rules
- LLMs do not learn the symbolic rules well
- Guiding them to learn improves their performance
- Future: more genetic ways to learn more software knowledge